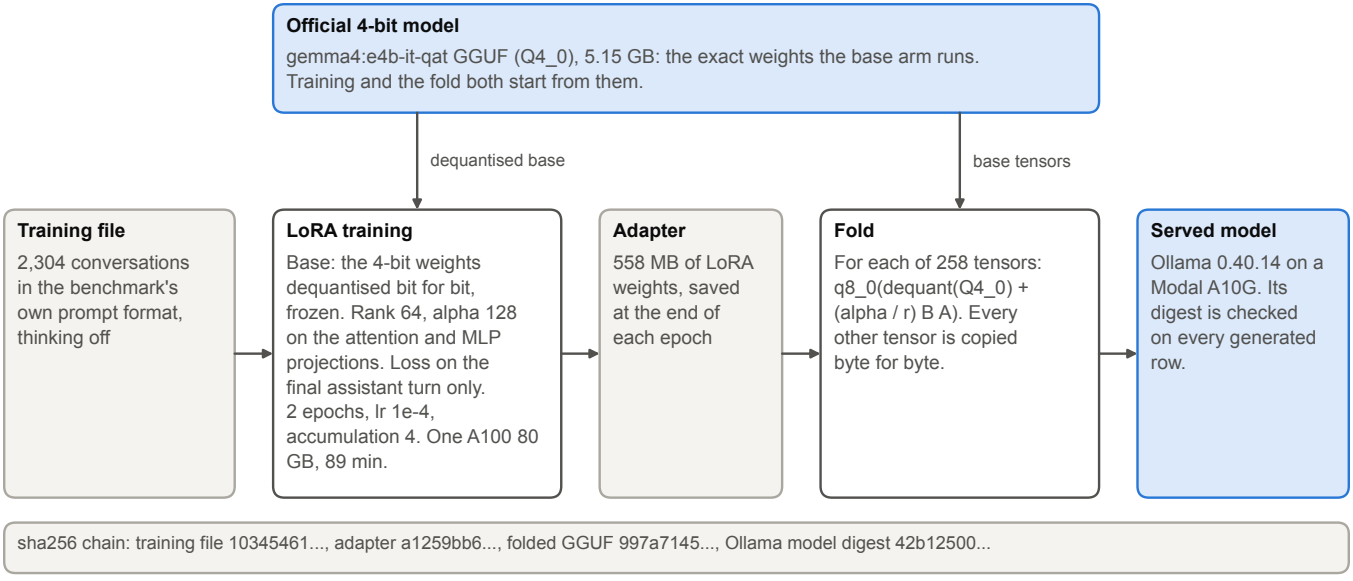
Official 4-bit model
gemma4:e4b-it-qat GGUF (Q4_0), 5.15 GB: the exact weights the base arm runs. Training and the fold both start from them.
dequantised base
base tensors
Training file
2,304 conversations in the benchmark's own prompt format, thinking off
LoRA training
Base: the 4-bit weights dequantised bit for bit, frozen. Rank 64, alpha 128 on the attention and MLP projections. Loss on the final assistant turn only. 2 epochs, lr 1e-4, accumulation 4. One A100 80 GB, 89 min.
Adapter
558 MB of LoRA weights, saved at the end of each epoch
Fold
For each of 258 tensors: q8_0(dequant(Q4_0) + (alpha / r) B A). Every other tensor is copied byte for byte.
Served model
Ollama 0.40.14 on a Modal A10G. Its digest is checked on every generated row.
sha256 chain: training file 10345461..., adapter a1259bb6..., folded GGUF 997a7145..., Ollama model digest 42b12500...