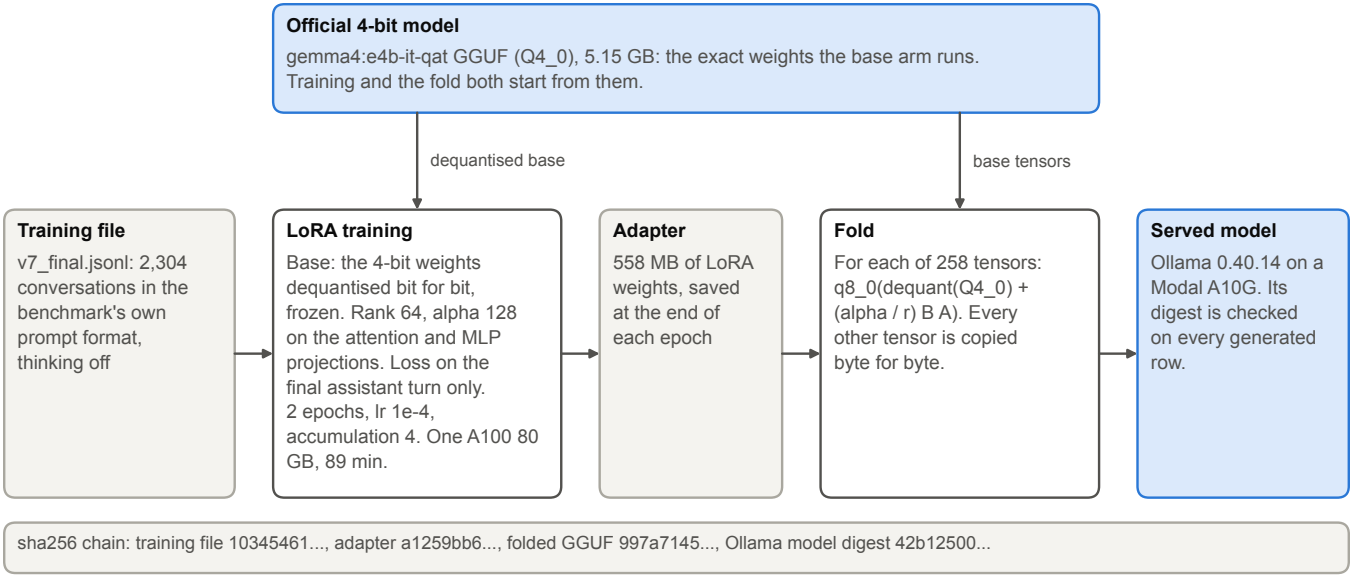
Official 4-bit model
gemma4:e4b-it-qat GGUF (Q4_0), 5.15 GB: the exact weights the base arm runs. Training and the fold both start from them.
dequantised base
base tensors
Training file
v7_final.jsonl: 2,304 conversations in the benchmark's own prompt format, thinking off
LoRA training
Base: the 4-bit weights dequantised bit for bit, frozen. Rank 64, alpha 128 on the attention and MLP projections. Loss on the final assistant turn only. 2 epochs, lr 1e-4, accumulation 4. One A100 80 GB, 89 min.
Adapter
558 MB of LoRA weights, saved at the end of each epoch
Fold
For each of 258 tensors: q8_0(dequant(Q4_0) + (alpha / r) B A). Every other tensor is copied byte for byte.
Served model
Ollama 0.40.14 on a Modal A10G. Its digest is checked on every generated row.
sha256 chain: training file 10345461..., adapter a1259bb6..., folded GGUF 997a7145..., Ollama model digest 42b12500...